

Aurimas Nausėdas

Vilnius, Lithuania

☎ +370 630 00386 • ☎ +44 7355 549475 • ✉ aurimas.nausedas@proton.me
[LinkedIn](#), [GitHub](#) - [aurimas13](#), [Portfolio](#) - [aurimas.io](#), [Weekly AI Newsletter](#)

AI engineer who builds LLM products end to end and the **evaluation layer** that decides whether they are allowed near a real decision: **agents** with explicit control loops and stop conditions, **RAG** pipelines whose answers are scored against criteria written in advance, guardrails, and per-call cost and latency telemetry. Three public agents built with no framework, each with its own evals, four AI prototypes live at [aurimas.io](#), and a regulated-finance background in AML and fraud detection. **Python since 2010; in AI engineering roles since May 2024; teaching Python and AI since 2022**, which is why I build with and explain to the people who will use the system.

Core Competencies

- **LLM and agent engineering:** Claude and OpenAI APIs; hand-written ReAct loops with tool calling, server-side tools and a bounded turn cap; streaming over Server-Sent Events; persistent memory and an answer cache in Supabase Postgres; human checkpoints where an action should not commit alone; MCP and reusable skills in daily use through **Claude Code**.
- **Evaluation, reliability and guardrails:** acceptance criteria and a pass threshold per criterion written before any run; each criterion scored separately, never averaged; a fixed evaluation set re-scored on a schedule to catch regressions; structured-output validation failure rate as a leading indicator; offline test suites on scripted fake model clients.
- **Retrieval, data and ML:** RAG pipelines with scored answers; SBERT and KeyBERT enrichment with entity de-duplication over Elasticsearch; OpenAI, Claude, Cohere and Gemini benchmarked on cost and quality; scikit-learn and PyTorch; SQL and ETL with data quality as a first-class concern.
- **Software engineering:** Python with FastAPI, Flask and Django; Go; TypeScript, React and Next.js; Node services with Stripe webhooks and cron jobs; REST APIs; pytest and ruff; Docker; Git; CI/CD with GitHub Actions; AWS; PostgreSQL and MS SQL Server.
- **Teaching, enablement and stakeholders:** two years teaching Python and AI part-time to hundreds of adults in Lithuania and Latvia; AI enablement across technical teams at Paysera; technical consulting at C&R Software; 100+ published editions written for readers who are not engineers.

Professional Experience

- **Vilniaus Asmodas** **Vilnius, Lithuania**
AI Product Engineer (Independent) *Oct 2024-Present*
 - Built and deployed three public agents with no framework, each with its own test suite in CI: [Code-Editing Agent](#) in Go, and [Claude Agent From Scratch](#) and [Calculator Agent](#) in Python (see Selected Work).
 - Built four AI prototypes end to end, owning each from problem statement and specification through build, evaluation and iteration, live at [aurimas.io](#): agentic AML investigation, LLM governance, PLG onboarding and a B2B agent workspace. Prototypes, not systems carrying production traffic.
 - Ran evaluation as a release gate: [Molecule to Machine](#) (100+ weekly editions) is assembled by reusable Claude Code skills that score each section against a written rubric until every criterion reaches 8 of 10; the skills write to a review folder and never publish, I do.
 - Kept the model out of steps that do not need one: for [100 Days With AI](#) a model writes and pushes each entry in the evening, while the 06:50 publishing run is deterministic, with no model in it.
- **Paysera** **Vilnius, Lithuania**
AI Lead *Aug 2024-Oct 2024*
 - Worked on AML and fraud-detection systems at an electronic money institution holding Bank of Lithuania Licence No. 1.
 - Helped colleagues adopt AI tools as everyday assistants and taught the technical teams: leadership by setting practice, not managing reports.
- **Royalty Range** **Remote**
AI / ML Engineer (contract) *May 2024-Jul 2024*
 - Defined the problem, built the scoring and flagging logic and shipped a transaction fraud-detection tool that flags suspicious transactions for human review, prioritising recall and analyst oversight over full automation; compared AI models for the task.

- **C&R Software** **Vilnius, Lithuania**
Technical Consultant *May 2023-Sep 2023*
 - Automated and optimised databases and workflows with AWS and Python; improved debt-collection ETL and data quality; built and tuned PostgreSQL and MS SQL Server databases.
- **Code Academy (LT) and LIKTA (LV)** **Lithuania and Latvia**
AI and Python Lecturer (part-time) *Jan 2022-Dec 2023*
 - Taught Python and AI to hundreds of adult learners across two countries, part-time; course materials public on GitHub ([CodeAcademy-AI-Course](#)).
- **Euromonitor; Thermo Fisher Scientific** **Vilnius, Lithuania**
Earlier technical roles *2022; 2016*
 - Euromonitor (Jan-Mar 2022): JavaScript web-scraping robots. Thermo Fisher Scientific (Jan-Mar 2016): HPLC and mass-spectrometry protein analysis, and scRNA-seq analysis with machine learning in Python.

Selected Work

- **Code-Editing Agent** ([code](#)): Go agent that reads, lists and edits files and researches the web, in a per-visitor sandbox; refuses ambiguous edits, redacts secrets before the model sees them, caps rounds, tokens, time and daily spend; every step streamed and stored. **39 deterministic evals on every commit** and **15/15 live-model cases**; its first day of live use surfaced 13 problems, each traced and fixed.
- **Claude Agent From Scratch** ([code](#)): ReAct agent with real tools, persistent memory, an answer cache and per-request cost accounting; 40 offline tests in CI. **Calculator Agent** ([code](#)): the smallest inspectable loop, 10-turn stop, ast allow-list instead of eval, 48 pytest tests.
- **ClearTrace Intelligence** ([code](#)): agentic AML investigation prototype; risk scoring, multi-hop transaction networks, SAR drafting, human-in-the-loop by design. **Aegis AI** ([code](#)): LLM governance prototype; every model call logged with tokens, cost and a PII check; a legacy-code to Python translator.
- **AI Gateway** ([code](#)): PLG onboarding prototype with funnel analytics, an A/B toggle and a streaming workspace of 16 agents. **Agentic Mobile** ([code](#)): B2B agent workspace with generative UI, approval cards and multi-step workflows.
- **Benchmarking AI Agent** (2024): screens comparable companies for transfer pricing from Elasticsearch and MongoDB data and returns accept or reject with reasons.
- **Open source and teaching**: [Machine-Learning-Goodness](#) (285+ stars), [Solutions-To-Problems](#) (94 stars, Python, Java, SQL), [100 Days With AI](#) (21+ stars).

Education

- **The University of Edinburgh** **Edinburgh, UK**
MChem, Master of Chemistry with Honours *2009-2015*
 - Integrated master's; research placements of 13 to 15 weeks each at Janelia (HHMI), LMU Munich (Amgen), EMBL and Nottingham, plus a University of Edinburgh placement across the 2013-2014 academic year.
- **Vilnius University** **Vilnius, Lithuania**
Doctoral studies, chemistry and bioinformatics (withdrawn) *Jan 2016-Dec 2021*
 - DNA-protein interaction analysis; co-developed a C++ analysis platform; passed doctoral exams.

Certifications

- **Anthropic (2026)**: Claude in Amazon Bedrock; Model Context Protocol: Advanced Topics; Introduction to MCP; Claude Code in Action; Claude Code 101; Introduction to Claude Cowork; Introduction to Agent Skills; AI Fluency for Builders; AI Fluency Framework and Foundations; Building with the Claude API (doing, expected - Oct 2026); Claude on Google Cloud (doing, expected - Oct 2026).
- **Microsoft (2024)**: Accelerating Materials Discovery with HPC and AI. **DeepLearning.AI (2021)**: Deep Learning Specialization. **Databricks (2023)**: Generative AI Fundamentals; Lakehouse Fundamentals (both expired Oct 2025).

Languages

- Lithuanian (native), English (C2), French (A1), Russian (A1, understanding only).

References

- Available upon request.